

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from nltk.tokenize import word_tokenize
from nltk.corpus import stopwords
from nltk.stem import WordNetLemmatizer
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import MultinomialNB, BernoulliNB # naive bayes classifier
from sklearn.linear_model import LogisticRegression
import sklearn.metrics as mt
from sklearn.feature_extraction.text import TfidfVectorizer
import plotly.express as px
import string
from wordcloud import WordCloud
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: data = pd.read_csv("https://raw.githubusercontent.com/rashakil-ds/Public-Datasets/m")
```

```
In [3]: data.head(5)
```

```
Out[3]:
```

	reviewText	Positive
0	This is a one of the best apps according to a b...	1
1	This is a pretty good version of the game for ...	1
2	this is a really cool game. there are a bunch ...	1
3	This is a silly game and can be frustrating, b...	1
4	This is a terrific game on any pad. Hrs of fun...	1

## Initial Observations

```
In [4]: df = data.copy()
```

```
In [5]: df.shape
```

```
Out[5]: (20000, 2)
```

```
In [6]: # We have 20000 rows and 2 columns
```

```
In [7]: df.nunique()
```

```
Out[7]: reviewText    20000
Positive              2
dtype: int64
```

```
In [8]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20000 entries, 0 to 19999
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  -
0   reviewText  20000 non-null  object
1   Positive    20000 non-null  int64
dtypes: int64(1), object(1)
memory usage: 312.6+ KB
```

```
In [9]: df.describe()
```

```
Out[9]:
```

	Positive
--	----------

count	20000.000000
-------	--------------

mean	0.761650
------	----------

std	0.426085
-----	----------

min	0.000000
-----	----------

25%	1.000000
-----	----------

50%	1.000000
-----	----------

75%	1.000000
-----	----------


max	1.000000
-----	----------

```
In [10]: df.dtypes
```

```
Out[10]: reviewText    object
Positive      int64
dtype: object
```

```
In [11]: label = ["Positive", "Negative"]
values = df["Positive"].value_counts()
plt.pie(labels = label, x = values)
```

```
Out[11]: ([<matplotlib.patches.Wedge at 0x27f2fe64e10>,
<matplotlib.patches.Wedge at 0x27f2ff00dd0>],
[Text(-0.8057580416543408, 0.748835080848488, 'Positive'),
Text(0.8057581117653374, -0.7488350054079725, 'Negative')])
```



We can clearly see class imbalance here

## Word Cloud

```
In [12]: positive_reviews = ' '.join(df[df['Positive'] == 1]['reviewText'])  
negative_reviews = ' '.join(df[df['Positive'] == 0]['reviewText'])
```

```
In [13]: wordcloud = WordCloud(width=800, height=400, background_color='white').generate(pos  
plt.figure(figsize=(10, 5))  
plt.imshow(wordcloud, interpolation='bilinear')  
plt.axis('off')  
plt.title("Positive Reviews")  
plt.show()
```

[illegible]

## Null Handling

```
In [15]: df.isna().sum()
```

```
Out[15]: reviewText      0  
Positive      0  
dtype: int64
```

## Text Preprocessing

```
In [16]: df["reviewText"][0]
```

```
Out[16]: 'This is a one of the best apps according to a bunch of people and I agree it has b  
ombs eggs pigs TNT king pigs and realistic stuff'
```

```
In [17]: stopwords.fileids()
```

```
Out[17]: ['arabic',  
          'azerbaijani',  
          'basque',  
          'bengali',  
          'catalan',  
          'chinese',  
          'danish',  
          'dutch',  
          'english',  
          'finnish',  
          'french',  
          'german',  
          'greek',  
          'hebrew',  
          'hinglish',  
          'hungarian',  
          'indonesian',  
          'italian',  
          'kazakh',  
          'nepali',  
          'norwegian',  
          'portuguese',  
          'romanian',  
          'russian',  
          'slovene',  
          'spanish',  
          'swedish',  
          'tajik',  
          'turkish']
```

```
In [18]: stopword = stopwords.words("english")
```

```
In [19]: punctuation = string.punctuation
```

```
In [20]: def preprocess_text(text):  
          preprocess_punctuation = "".join([char for char in text if char not in punctuat
```

```
preprocess_stopword = [char for char in preprocess_punctuation.split() if char.  
return preprocess_stopword
```

```
In [21]: df["reviewText"] = df["reviewText"].apply(preprocess_text)
```

```
In [22]: df.head()
```

```
Out[22]:
```

	reviewText	Positive
0	[one, best, apps, acording, bunch, people, agr...	1
1	[pretty, good, version, game, free, LOTS, diff...	1
2	[really, cool, game, bunch, levels, find, gold...	1
3	[silly, game, frustrating, lots, fun, definite...	1
4	[terrific, game, pad, Hrs, fun, grandkids, lov...	1

## Text Normalization/Scaling

```
In [23]: lemmatizer = WordNetLemmatizer()
```

```
In [24]: df["reviewText"][0]
```

```
Out[24]: ['one',  
          'best',  
          'apps',  
          'acording',  
          'bunch',  
          'people',  
          'agree',  
          'bombs',  
          'eggs',  
          'pigs',  
          'TNT',  
          'king',  
          'pigs',  
          'realustic',  
          'stuff']
```

```
In [25]: def normalize_data(text):  
          normalized_text = ' '.join([lemmatizer.lemmatize(word) for word in text])  
          return normalized_text
```

```
In [26]: import nltk #Handle Lookup Error  
          nltk.download("wordnet")  
          !unzip /usr/share/nltk_data/corpora/wordnet.zip -d /usr/share/nltk_data/corpora/
```

```
[nltk_data] Downloading package wordnet to
[nltk_data] C:\Users\ANJAR\AppData\Roaming\nltk_data...
[nltk_data] Package wordnet is already up-to-date!
unzip: cannot find or open /usr/share/nltk_data/corpora/wordnet.zip, /usr/share/nltk_data/corpora/wordnet.zip.zip or /usr/share/nltk_data/corpora/wordnet.zip.ZIP.
```

```
In [27]: df["reviewText"] = df["reviewText"].apply(normalize_data)
```

```
In [28]: df.head()
```

```
Out[28]:
```

	reviewText	Positive
0	one best apps acording bunch people agree bomb...	1
1	pretty good version game free LOTS different l...	1
2	really cool game bunch level find golden egg s...	1
3	silly game frustrating lot fun definitely reco...	1
4	terrific game pad Hrs fun grandkids love Great...	1

## Vectorization

```
In [29]: vectorizer = TfidfVectorizer()
```

```
In [30]: X = vectorizer.fit_transform(df["reviewText"])
y = df["Positive"]
```

```
In [31]: X.shape
```

```
Out[31]: (20000, 22617)
```

## Data Splitting

```
In [32]: xtrain, xtest, ytrain, ytest = train_test_split(X, y, test_size=0.3, random_state=4)
```

```
In [33]: xtrain.shape
```

```
Out[33]: (14000, 22617)
```

## Model Training

```
In [34]: mnbNB = MultinomialNB()
berNB = BernoulliNB()
lr = LogisticRegression()
```

```
In [35]: models = [mnbnb, bernb, lr]
```

```
In [36]: def train_models(X_train,y_train, models):  
    for model in models:  
        model.fit(X_train,y_train)  
        print(f"{model} : {mt.accuracy_score(y_train, model.predict(X_train))}")
```

```
In [37]: train_models(xtrain,ytrain,models)
```

```
MultinomialNB() : 0.8077857142857143  
BernoulliNB() : 0.9132857142857143  
LogisticRegression() : 0.9223571428571429
```

## Model Evaluation

```
In [38]: def evaluate_model(X_test, y_test, models):  
  
    y_predicted_list = []  
    accuracy_scores = []  
  
    for model in models:  
        prediction = model.predict(X_test)  
        y_predicted_list.append(prediction)  
        accuracy_scores.append(mt.accuracy_score(y_test,prediction))  
  
    for idx, prediction in enumerate(y_predicted_list):  
        print(f"{models[idx]}")  
        print()  
        print("Test Accuracy Score: ",accuracy_scores[idx])  
        print()  
        print(mt.classification_report(y_test,prediction))  
        print("-----")  
  
    return accuracy_scores, y_predicted_list
```

```
In [39]: accuracy_scores, predictions = evaluate_model(xtest,ytest,models)
```



```
MultinomialNB()
```

```
Test Accuracy Score: 0.7873333333333333
```

	precision	recall	f1-score	support
0	0.91	0.11	0.19	1411
1	0.78	1.00	0.88	4589
accuracy			0.79	6000
macro avg	0.85	0.55	0.53	6000
weighted avg	0.81	0.79	0.72	6000

```
-----  
BernoulliNB()
```

```
Test Accuracy Score: 0.8635
```

	precision	recall	f1-score	support
0	0.80	0.56	0.66	1411
1	0.88	0.96	0.91	4589
accuracy			0.86	6000
macro avg	0.84	0.76	0.79	6000
weighted avg	0.86	0.86	0.85	6000

```
-----  
LogisticRegression()
```

```
Test Accuracy Score: 0.8843333333333333
```

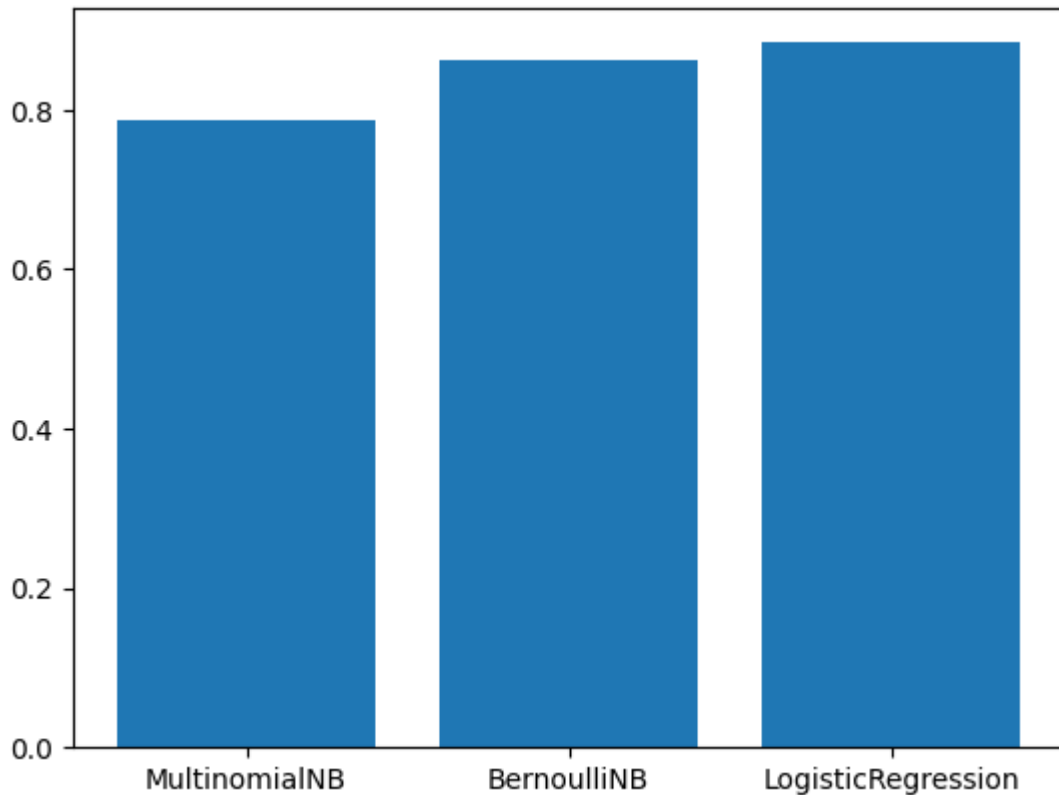
	precision	recall	f1-score	support
0	0.86	0.60	0.71	1411
1	0.89	0.97	0.93	4589
accuracy			0.88	6000
macro avg	0.88	0.79	0.82	6000
weighted avg	0.88	0.88	0.88	6000

-----  
By analyzing the classification reports above, we can see that Logistic Regression has the best accuracy with 88%.

## Accuracy Comparison

```
In [40]: plt.bar(['MultinomialNB', 'BernoulliNB', 'LogisticRegression'], accuracy_scores)
```

```
Out[40]: <BarContainer object of 3 artists>
```



## Multinomial Naive Bayes:

- **Test Accuracy Score:** 78.73%
- **Precision-Recall-F1 Score:**
  - Class 0: Precision 91%, Recall 11%, F1-Score 19%
  - Class 1: Precision 78%, Recall 100%, F1-Score 88%

An interesting start with a good overall accuracy, though there's room for improvement in predicting Class 0.

## Bernoulli Naive Bayes:

- **Test Accuracy Score:** 86.35%
- **Precision-Recall-F1 Score:**
  - Class 0: Precision 80%, Recall 56%, F1-Score 66%
  - Class 1: Precision 88%, Recall 96%, F1-Score 91%

Bernoulli Naive Bayes exhibits strong performance, particularly in correctly identifying instances of Class 1.

## Logistic Regression:

- **Test Accuracy Score:** 88.43%

- **Precision-Recall-F1 Score:**
  - Class 0: Precision 86%, Recall 60%, F1-Score 71%
  - Class 1: Precision 89%, Recall 97%, F1-Score 93%

Logistic Regression takes the lead with the highest accuracy and robust performance across both classes.

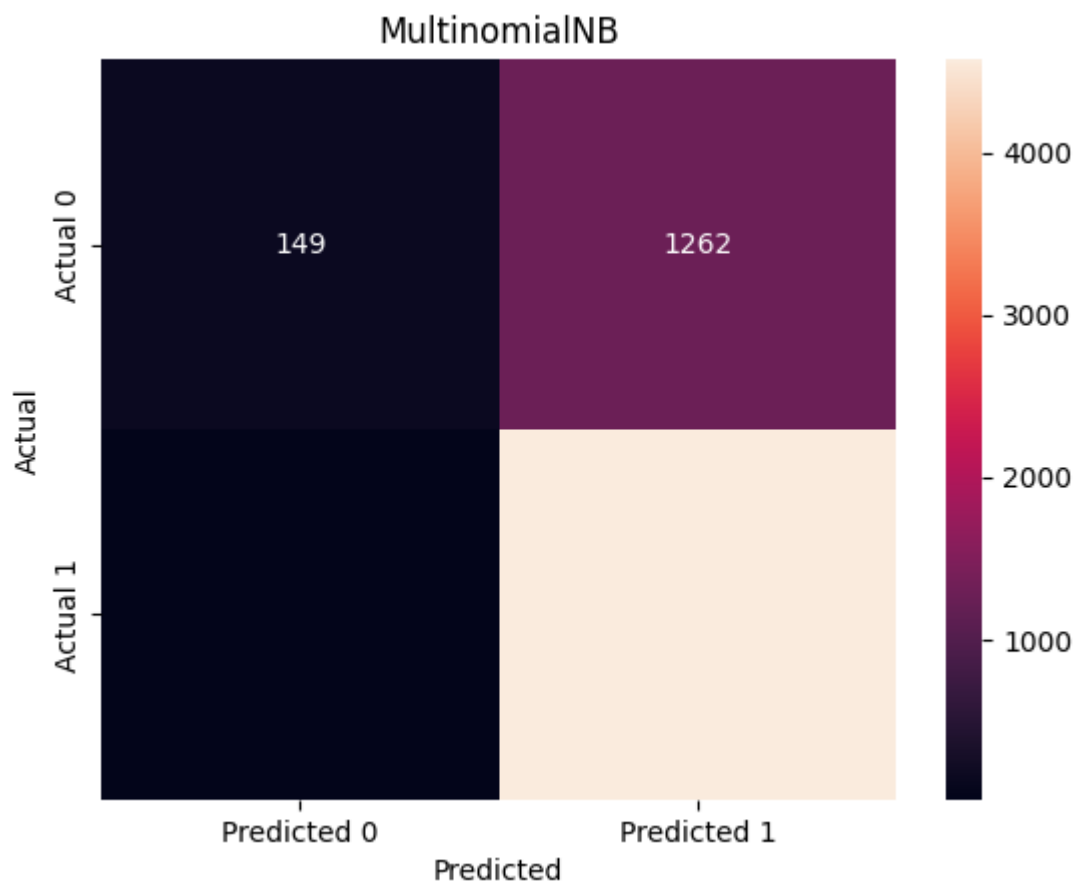
## Insights:

- Multinomial Naive Bayes shows potential but may benefit from additional optimization, especially in predicting Class 0.
- Bernoulli Naive Bayes excels in predicting Class 1, indicating its effectiveness in capturing relevant patterns.
- Logistic Regression emerges as a strong contender, offering high accuracy and balanced performance.

## Confusion Matrix

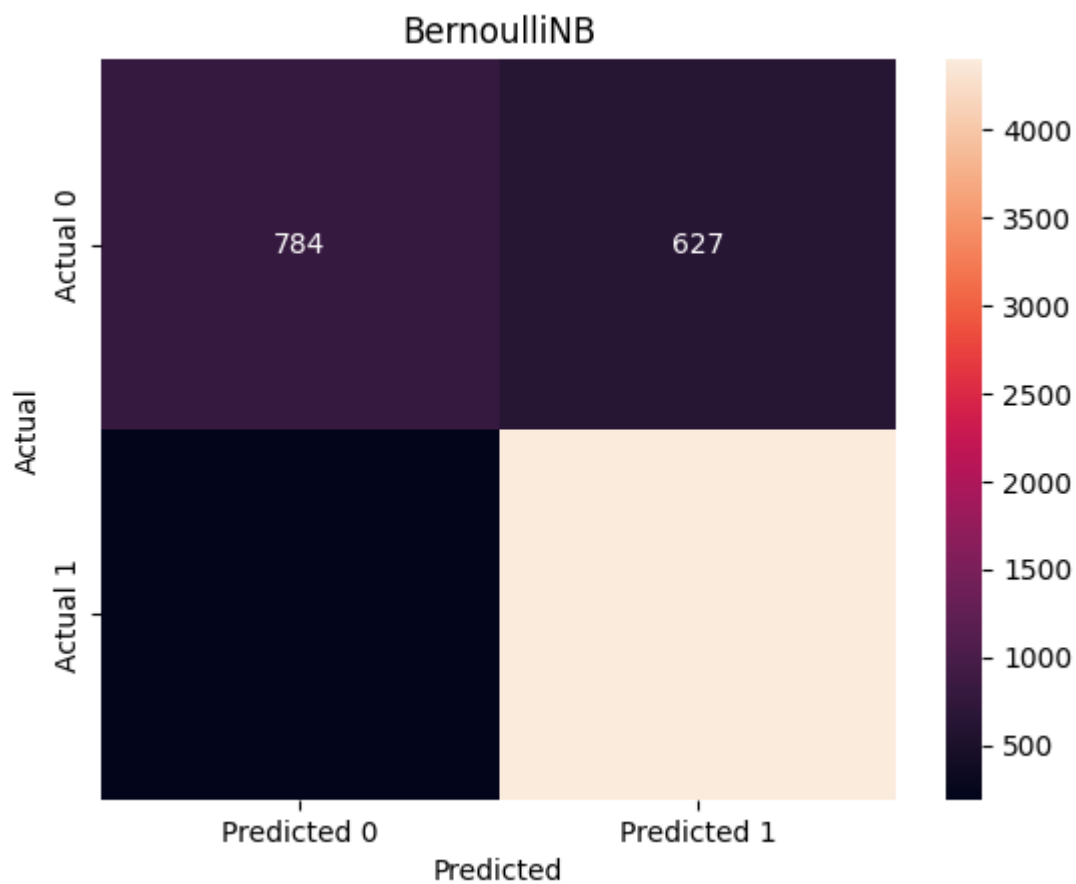
```
In [41]: sns.heatmap(mt.confusion_matrix(ytest, predictions[0]),annot= True,fmt="d",yticklab
plt.title("MultinomialNB")
plt.xlabel('Predicted')
plt.ylabel('Actual')
```

```
Out[41]: Text(50.72222222222214, 0.5, 'Actual')
```




```
In [42]: sns.heatmap(mt.confusion_matrix(ytest, predictions[1]),annot= True,fmt="d",yticklab  
plt.title("BernoulliNB")  
plt.xlabel('Predicted')  
plt.ylabel('Actual')
```

```
Out[42]: Text(50.72222222222214, 0.5, 'Actual')
```



```
In [43]: sns.heatmap(mt.confusion_matrix(ytest, predictions[2]),annot= True,fmt="d",yticklab  
plt.title("LogisticRegression")  
plt.xlabel('Predicted')  
plt.ylabel('Actual')
```

```
Out[43]: Text(50.72222222222214, 0.5, 'Actual')
```



## Multinomial Naive Bayes:

- **True Positives (TP):** 4575
- **True Negatives (TN):** 149
- **False Positives (FP):** 1262
- **False Negatives (FN):** 14

## Observations:

- The model performs well in correctly predicting Class 1 (heart attack occurrence) with a high True Positive count.
- However, it struggles in predicting instances of Class 0, as indicated by the low True Negative count and a relatively high False Positive count.

## Bernoulli Naive Bayes:

- **True Positives (TP):** 4397
- **True Negatives (TN):** 784
- **False Positives (FP):** 627
- **False Negatives (FN):** 192

## Observations:

- The model shows a strong ability to predict both Class 0 and Class 1, with high counts in both True Positives and True Negatives.
- The False Positive count is relatively low, indicating a good balance between precision and recall.

## Logistic Regression:

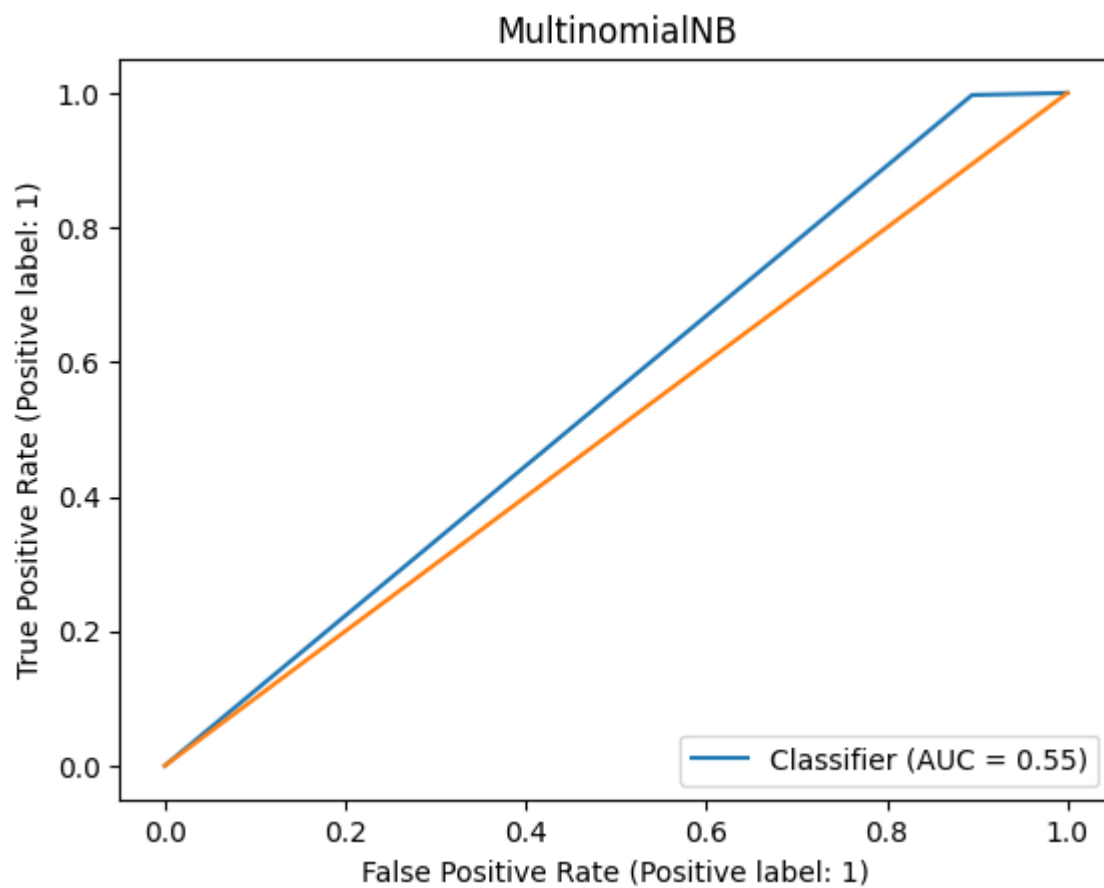
- **True Positives (TP):** 4454
- **True Negatives (TN):** 852
- **False Positives (FP):** 559
- **False Negatives (FN):** 135

## Observations:

- Logistic Regression demonstrates a balanced performance in predicting both classes, with high counts in both True Positives and True Negatives.
- The False Positive count is relatively low, contributing to the model's high precision and accuracy.

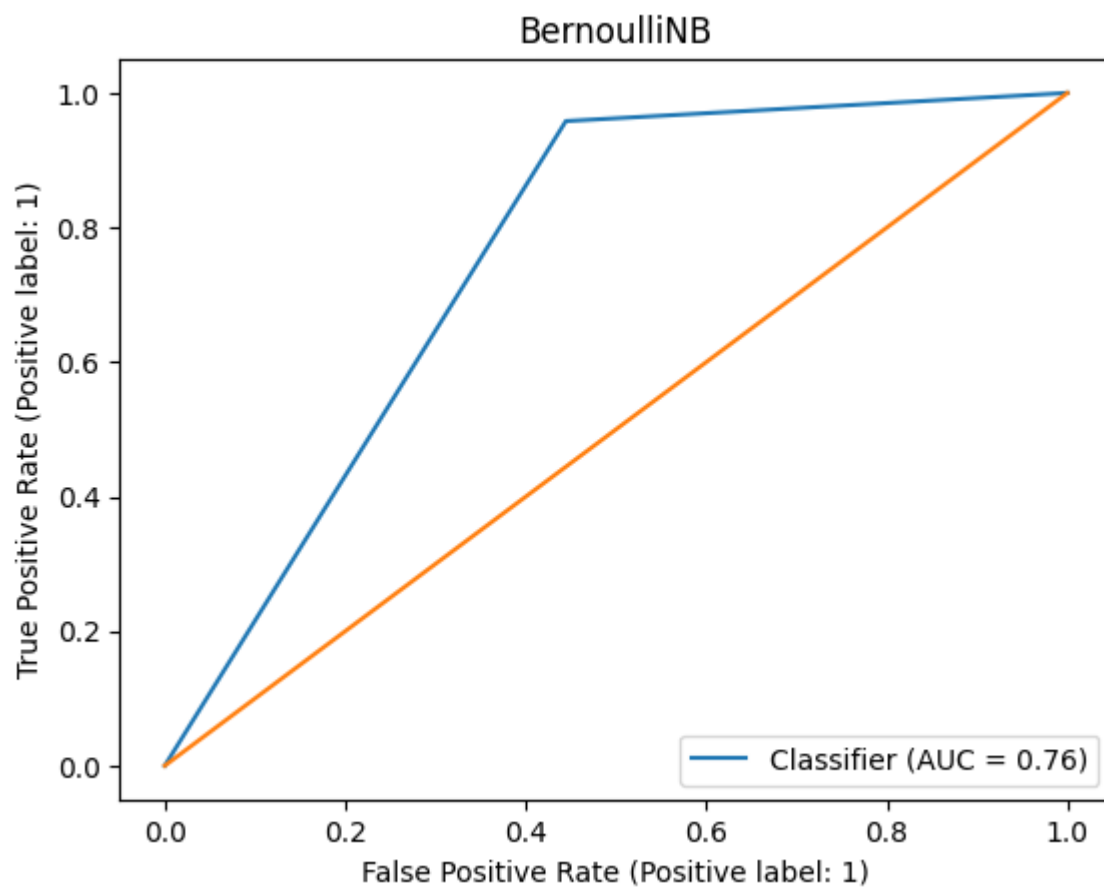
## Roc Auc Curve

```
In [44]: mt.RocCurveDisplay.from_predictions(ytest, predictions[0])
plt.plot([0,1],[0,1])
plt.title("MultinomialNB")
plt.show()
```




```
In [45]: mt.RocCurveDisplay.from_predictions(ytest, predictions[1])
plt.plot([0,1],[0,1])
plt.title("BernoulliNB")
plt.show()
```





```
In [46]: mt.RocCurveDisplay.from_predictions(ytest, predictions[2])
plt.plot([0,1],[0,1])
plt.title("LogisticRegression")
plt.show()
```



The Logistic Regression Model performs better compared to other two models in ROC AUC Curve with a coverage of 79% percent.

## HyperParameter Tuning

*As Logistic Regression performed the best across all models, we hypertune the parameters to potentially gain even more performance*

```
In [47]: from sklearn.model_selection import GridSearchCV
```

```
In [48]: param_grid = {
    'penalty': ['l1', 'l2'],           # Regularization type (L1 or L2)
    'C': [0.001, 0.01, 0.1, 1, 10],  # Inverse of regularization strength
    'solver': ['liblinear', 'saga']   # Algorithm to use for optimization
}
```

```
In [49]: scoring = {
    'accuracy': mt.make_scorer(mt.accuracy_score),
    'precision': mt.make_scorer(mt.precision_score, pos_label='positive'),
    'recall': mt.make_scorer(mt.recall_score, pos_label='positive'),
    'f1': mt.make_scorer(mt.f1_score, pos_label='positive')
}
```

```
In [50]: grid_search = GridSearchCV(estimator=lr, param_grid=param_grid, cv=5, scoring=scori
grid_search.fit(xtrain, ytrain)
```

```
Out[50]: ▸ GridSearchCV
▸ estimator: LogisticRegression
    ▸ LogisticRegression
```

```
In [51]: best_params = grid_search.best_params_
```

```
In [52]: lr_best = LogisticRegression(**best_params)
```

```
In [53]: lr_best.fit(xtrain,ytrain)
```

```
Out[53]: ▾ LogisticRegression
LogisticRegression(C=10, solver='liblinear')
```

```
In [54]: predicted = lr_best.predict(xtest)
```

```
In [55]: print(mt.classification_report(ytest,predicted))
```

	precision	recall	f1-score	support
0	0.81	0.71	0.76	1411
1	0.91	0.95	0.93	4589
accuracy			0.89	6000
macro avg	0.86	0.83	0.84	6000
weighted avg	0.89	0.89	0.89	6000

## Summary:

- All three models perform well in predicting Class 1 (heart attack occurrence), with high True Positive counts.
- Multinomial Naive Bayes struggles more with predicting instances of Class 0, while Bernoulli Naive Bayes and Logistic Regression demonstrate better balance in predicting both classes.
- Logistic Regression stands out with the highest overall accuracy and balanced performance across both cl
- After using Hyperparameter tuning, Logistic regression gained slight performance increase of about 1%ass

## Recommendations:

- For Multinomial Naive Bayes, consider addressing the imbalance in predicting Class 0 through further optimization or sampling techniques.

In [ ]: